

VE482 Lab Report

Lab 7 - Fall 2020

Wu Qinhang (518370910041)

Sun Zhimin (518370910219)

Table of Contents

- Linux kernel source code
- clean software updates

Kernel Module

dadfs

mutex

information share

changes

Reference

Kernel Module

What is a kernel module, and how does it differ from a regular library? [1](#) [2](#) [3](#)

- A kernel module is an object file aimed at extending the kernel functionality when it is loaded onto the kernel. It can be treated as a plugin to the kernel program. A kernel module is linked only to the kernel, and it can be dynamically loaded. In this way, when a new kernel module is introduced, the kernel doesn't need to be re-compiled. However, a regular library is just a collection of relevant functionality that supports different application. If a library of kernel is updated, the kernel needs to be re-compiled.

How to compile a kernel module?

- The module should not be linked to libraries, and the module should be compiled with `kbuild`.

The following example is taken from [1](#):

```
1 // helloworld.c
2 #include <linux/kernel.h>
3 #include <linux/init.h>
4 #include <linux/module.h>
5
6 MODULE_DESCRIPTION("My kernel module");
7 MODULE_AUTHOR("Me");
8 MODULE_LICENSE("GPL");
9
10 static int dummy_init(void)
11 {
12     pr_debug("Hi\n");
13     return 0;
```

```

14 }
15
16 static void dummy_exit(void)
17 {
18     pr_debug("Bye\n");
19 }
20
21 module_init(dummy_init);
22 module_exit(dummy_exit);

```

```

1 # Makefile
2 KDIR = /lib/modules/`uname -r`/build
3
4 kbuild:
5     make -C $(KDIR) M=`pwd`
6
7 clean:
8     make -C $(KDIR) M=`pwd` clean

```

Kbuild

```

1 EXTRA_CFLAGS = -Wall -g
2 obj-m          = helloworld.o

```

dadfs

For source code, see [./dadfs/*](#).

Build command:

```

1 #!/bin/bash
2 # check the magic version
3 uname -r
4 # check the loop status
5 lsblk --fs
6 make
7 dd bs=4096 count=100 if=/dev/zero of=disk
8 dd bs=1M count=10 if=/dev/zero of=journal
9 mke2fs -b 4096 -o journal_dev journal
10 ./mkfs-dadfs disk
11 # insert the module
12 sudo insmod ./dadfs.ko
13 test_journal_dev=$(sudo losetup -f --show journal)
14 sudo mount -o loop,owner,group,users,journal_path="$test_journal_dev" -t
   dadfs disk /mnt

```

```

william@ubuntu:~/code/ve482-17/VE482-Lab7$ sudo chmod 777 /mnt/awordfromdad
william@ubuntu:~/code/ve482-17/VE482-Lab7$ sudo ls -al /mnt
total 0
-rwxrwxrwx 1 root root 0 Nov 18 06:26 awordfromdad
william@ubuntu:~/code/ve482-17/VE482-Lab7$ cat /mnt/awordfromdad
cat: /mnt/awordfromdad: Permission denied
william@ubuntu:~/code/ve482-17/VE482-Lab7$ sudo cat /mnt/awordfromdad
Congratulations, I'm proud of you. Dad

```

Figure 1. "I'm proud of you, Dad"

mutex

How are mutex defined and used? How good is this approach?

Three mutex in total are defined, they are:

- `darfs_sb_lock` Used for locking super block when action is taken on block
- `darfs_inodes_mgmt_lock` Used for locking super block when block inode need to be changed
- `darfs_directory_children_update_lock` Used for locking super block when children of the block is changed

Advantages:

1. Naming of variables are pretty clear.
2. Function can be accomplished with one mutex, but it separates into three parts in order to improve the efficiency.
3. Of course dad your code is perfect! No one knows mutex BETTER then you!

information share

How is information shared between the kernel and user spaces?

First, we need to let kernel allocate the memory, then allow the userspace process to map it into its address space with `mmap()`.

Then, take C as an example, the kernel thread and the userspace process will use different pointers to refer to the memory region:

- Kernel thread will use a pointer to the memory allocated within the kernel address space,
- User space process will use a pointer to the memory region returned by `mmap()`.

changes

```

1  11c11
2  < #include <linux/slab.h>
3  ---
4  > #include <linux/slab.h>
5  17,19d16
6  < #if LINUX_VERSION_CODE >= KERNEL_VERSION(3, 11, 0)
7  < #include <linux/uio.h>
8  < #endif
9  358d354
10 < #if LINUX_VERSION_CODE < KERNEL_VERSION(4, 0, 0)
11 361,363d356

```

```

12 < #else
13 < ssize_t dadfs_write(struct kiocb * kiocb, struct iov_iter * iov_iter)
14 < #endif
15 373d365
16 <         // struct *i_mapping;
17 379c371
18 < #if LINUX_VERSION_CODE < KERNEL_VERSION(4, 0, 0)
19 ---
20 >
21 381,383d372
22 < #else
23 <         sb = kiocb->ki_filp->f_inode->i_sb;
24 < #endif
25 389,390d377
26 <
27 < #if LINUX_VERSION_CODE < KERNEL_VERSION(4, 0, 0)
28 392,396d378
29 <         // int generic_write_checks(struct file *file, loff_t *pos,
30         size_t *count, int isblk);
31 < #else
32 <         retval = generic_write_checks(kiocb, iov_iter);
33 <         // extern ssize_t generic_write_checks(struct kiocb *, struct
34         iov_iter *);
35 < #endif
36 399c381
37 < #if LINUX_VERSION_CODE < KERNEL_VERSION(4, 0, 0)
38 ---
39 >
40 401,404d382
41 < #else
42 <         inode = kiocb->ki_filp->f_inode;
43 < #endif
44 <
45 406c384
46 < #if LINUX_VERSION_CODE < KERNEL_VERSION(4, 0, 0)
47 ---
48 >
49 409,412d386
50 < #else
51 <         bh = sb_bread(kiocb->ki_filp->f_inode->i_sb,
52         sfs_inode-
53         >data_block_number);
54 < #endif
55 422d395
56 < #if LINUX_VERSION_CODE < KERNEL_VERSION(4, 0, 0)
57 424,426d396
58 < #else
59 <         buffer += kiocb->ki_pos;
60 < #endif
61 434,440c404,406
62 < #if LINUX_VERSION_CODE < KERNEL_VERSION(4, 0, 0)

```

```

60 <         if (copy_from_user(buffer, buf, len)) { // (void *to, const
        void __user *from, __kernel_size_t n);
61 < #else
62 <         if (copy_from_iter(buffer, iov_iter->count, iov_iter) == 0) {
        // FIXME
63 <             // size_t copy_from_iter(void *addr, size_t bytes,
        struct iov_iter *i)
64 < #endif
65 <             brelse(bh);
66 ---
67 >
68 >         if (copy_from_user(buffer, buf, len)) {
69 >             brelse(bh);
70 445d410
71 < #if LINUX_VERSION_CODE < KERNEL_VERSION(4, 0, 0)
72 447,449d411
73 < #else
74 <         kiocb->ki_pos += iov_iter->count; // FIXME: len
75 < #endif
76 477,478d438
77 <
78 < #if LINUX_VERSION_CODE < KERNEL_VERSION(4, 0, 0)
79 480,482d439
80 < #else
81 <         sfs_inode->file_size = (kiocb->ki_pos);
82 < #endif
83 485d441
84 < #if LINUX_VERSION_CODE < KERNEL_VERSION(4, 0, 0)
85 488a445
86 >
87 490,497d446
88 < #else
89 <         ;
90 <         mutex_unlock(&dadfs_inodes_mgmt_lock);
91 <         return retval;
92 <     }
93 <     mutex_unlock(&dadfs_inodes_mgmt_lock);
94 <     return iov_iter->count; // TODO: change len
95 < #endif
96 502d450
97 < #if LINUX_VERSION_CODE < KERNEL_VERSION(4, 0, 0)
98 504,506d451
99 < #else
100 <     .write_iter = dadfs_write,
101 < #endif
102 664,666d608
103 <
104 < // dadfs_bmap
105 < // TODO: address_space_operations dadfs_a_ops = {}

```

Reference

1. "Kernel modules — The Linux Kernel documentation," *Github.io* , 2020. [https://linux-kernel-labs.github.io/refs/heads/master/labs/kernel_modules.html#:~:text=A%20kernel%20module%20\(or%20loadable,the%20form%20of%20kernel%20modules](https://linux-kernel-labs.github.io/refs/heads/master/labs/kernel_modules.html#:~:text=A%20kernel%20module%20(or%20loadable,the%20form%20of%20kernel%20modules). (accessed Nov. 13, 2020). [↵](#) [↵](#)
2. jetru, "Difference between a module, library and a framework," *Stack Overflow* , Nov. 04, 2010. <https://stackoverflow.com/questions/4099975/difference-between-a-module-library-and-a-framework> (accessed Nov. 13, 2020). [↵](#)
3. *Oracle.com* , 2010. <https://docs.oracle.com/cd/E19253-01/817-5789/emjlr/index.html> (accessed Nov. 14, 2020). [↵](#)