

VE482 Lab Report

Lab 11 - Fall 2020

Boming Zhang

Chujie Ni

Qinhang Wu

Zhimin Sun

Goals of the lab

- Understand filesystems
- Work with a filesystem in userspace
- Learn the basics of FUSE

VE482 Lab 11

What is a filesystem?

How is the Linux VFS working?

What is FUSE, and how does it interact with the VFS? Can you sketch it quickly to make it clearer?

Implementation

Task done

Demo

Reference

VE482 Lab 11

What is a filesystem?

File system is a system that controls how data is stored and retrieved. By separating the data into pieces and giving each piece a name, the data is easily isolated and identified. Taking its name from the way paper-based data management system is named, each group of data is called a "file." The structure and logic rules used to manage the groups of data and their names is called a "file system."

How is the Linux VFS working?

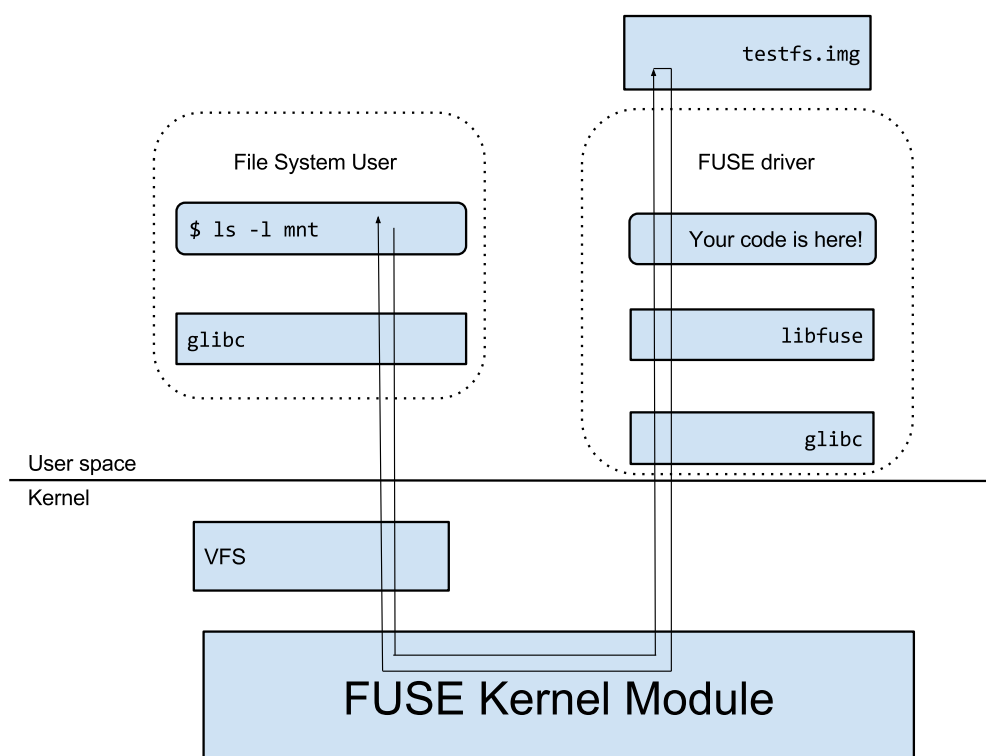
VFS is the software layer in the kernel that provides the filesystem interface to userspace programs. It also provides an abstraction within the kernel which allows different filesystem implementations to coexist.

The key point is that VFS is a layer of code which implements generic filesystem actions and vectors requests to the correct specific code to handle the request. A VFS specifies an interface between the kernel and a concrete file system. Therefore, it is easy to add support for new file system types to the kernel simply by fulfilling the API. Therefore, we can only focus on the API provided by VFS instead of all the details associated with the concrete file system.

What is FUSE, and how does it interact with the VFS? Can you sketch it quickly to make it clearer?

FUSE for Filesystem in User Space. An open source framework for implementing filesystem in user land.

FUSE has a daemon process that deal with the file operations. For every file operations, it take these operations in user space and convert it into kernel space commands and finish the job, twice each operation, one for request and the other for response.



Implementation

We first try to recover the file. After download `del.tar.bz2`, run:

```
1 # cd .del
2 # vim -r lemondbfs.c
```

And we will get the recovered version of `lemondbfs.c`.

Then we implement the lemonDBfs in `./lemondbfs/`.

Task done

1. for each .tbl file show a dir in mount_dir
eg. for file mTable10.tbl show dir mtable10
2. inside each dir show the corresponding table in file orig.tbl
eg. for file mTable10.tbl show mtable10/orig.tbl
3. inside each dir show an empty file called .query
eg. for mTable10.tbl show the empty file mTable10/.query
4. when a query is written to a .query file, apply it in the orig.tbl and create the file res.tbl containing the result of the query
eg. upon a write on mTable10/.query, run the query (error on wrong query) and create mTable10/res.tbl with the result of the query
5. lemonDBfs becomes more stable (not allow to run as root)
6. Allow more than one query in the .query file
7. When many queries are input allow temporary result tables to be created.
8. Allow the deletion of .tbl files: when a dir is deleted in mount_dir the corresponding .tbl file should disappear

Demo

```
1 $ cd /lemondbfs/src
2 $ make
3 (so many outputs)
4 $ ls ../db_dir/
5 mTable0.tbl mTable1.tbl mTable3.tbl
6 $ ./lemondbfs ../db_dir/ ../mount_dir/
7 about to call fuse_main
8 lm_data: 0x563217b8deb0
9 $ ls ../mount_dir/
10 mTable0 mTable1 mTable3
11 $ ls -lha ../mount_dir/mTable0/
12 total 0
13 drwxr-xr-x 1 root root  0 Jan  1 1970 .
14 drwxr-xr-x 2 root root  0 Jan  1 1970 ..
15 -rwxr-xr-x 1 root root 65K Jan  1 1970 orig.tbl
16 -rwxr-xr-x 1 root root  0 Jan  1 1970 .query
17 $ cat ../mount_dir/mTable0/orig.tbl
18 (so many outputs)
```

```
19 $ echo "UPDATE ( c0 100 ) FROM mTable0 WHERE ( c2 > 0 );" >
    ../mount_dir/mTable0/.query
20 $ ls -lha ../mount_dir/mTable0/
21 total 0
22 drwxr-xr-x 1 root root  0 Jan  1 1970 .
23 drwxr-xr-x 2 root root  0 Jan  1 1970 ..
24 -rwxr-xr-x 1 root root 65K Jan  1 1970 orig.tbl
25 -rwxr-xr-x 1 root root  0 Jan  1 1970 .query
26 -rwxr-xr-x 1 root root 104K Jan  1 1970 res.
27 $ rmdir ../mount_dir/mTable3/
28 $ ls ../mount_dir/
29 mTable0  mTable1
30 $ ls ../db_dir/
31 mTable0.tbl  mTable1.tbl
```

Reference

- https://en.wikipedia.org/wiki/File_system
- https://en.wikipedia.org/wiki/Virtual_file_system
- <https://ic.unicamp.br/~islene/2s2013-mo806/vfs/andre-zhen.pdf>
- <https://sites.cs.ucsb.edu/~trinabh/classes/w19/labs/lab6.html>