


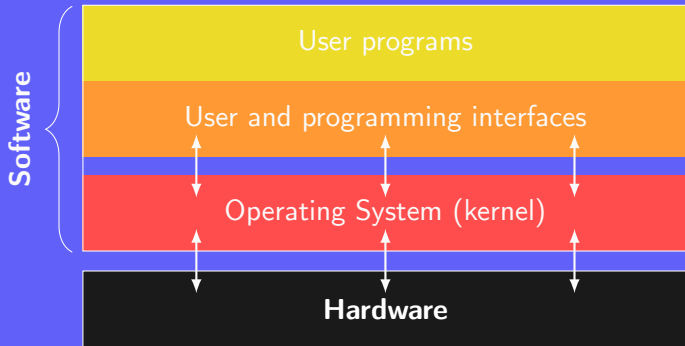


Introduction to Operating Systems

1. Operating systems overview

Manuel – Fall 2020







Job of an Operating System (OS):

- Manage and assign the hardware resources
- Hide complicated details to the end user
- Provide abstractions to ease interactions with the hardware

The early days:

- Birth of modern computing: 19th century (Babbage)
- Vacuum tube: 1945–1955 (1st generation)
- Transistor: 1955–1965 (2nd generation)





Using the device:

- Program at most 40 steps
- Read input from cardboards
- Wire them on a plugboard
- Punch output on cardboards

Multiprogramming: 1965–1980 (3rd generation)

- Multiple jobs kept in memory at the same time
- CPU multiplexed among them

Multiprogramming: 1965–1980 (3rd generation)

- Multiple jobs kept in memory at the same time
- CPU multiplexed among them

Multiprogramming requires:

- Memory management: allocate memory to several jobs
- CPU scheduling: choose a job to be run
- Simultaneous Peripheral Operation On Line (SPOOL): load a new job from disk, run it, output it on disk

Most famous OS:


- Disk Operating System (DOS)
- DOS/Basic package sold to computer companies
- MS-DOS, including many features from UNIX
- GUI invented in the 1960s, then copied by Apple
- Microsoft copied Apple (Windows working on top of MS-DOS)
- Many OS derived from UNIX (MINIX, LINUX, BSD...)

Device and task oriented OS types:

- Personal Computers (PC)
- Servers: serve users over a network (print, web, IM...) → Solaris, FreeBSD, Linux, Windows Server
- Multi processors: multiple CPU in a single system → Linux, Windows, OS X...
- Handheld computers: PDA, smartphone
- Embedded devices: TV, microwave, DVD player, mp3 player, old cell phones → everything stored in ROM, much more simple OS

More device and task oriented OS types:

- Real-Time: time is key parameter (e.g. assembly line, army, avionics...) → overlap with embedded/handheld systems
- Mainframe: room-sized computers, data centers → OS oriented toward processing many jobs at once and efficient I/O
- Sensor node: tiny computers communicating between each other and a base station (guard border, intrusion/fire detection etc...). Composed of CPU RAM ROM (+other sensors), small battery → simple OS design TinyOS
- Smart card: credit card size with a CPU chip, severe memory/processing constraints → smallest/primitive OS



A computer is often composed of:

- CPU
- Memory
- Monitor + video controller
- Keyboard + keyboard controller
- Hard Disk Drive (HDD) + hard disk controller
- Bus

What are the controllers, and the bus?

Basics:

- CPU is the “computer’s brain”
- CPU can only execute a specific set of instructions
- CPU fetches instructions from the memory and executes them

Registers:

- General register: hold variables/temporary results
e.g. program counter: address of next instruction to fetch
- Stack pointer: parameters/variables not kept in registers
- Program Status Word (PSW): control bits





Superscalar:

- Multiple execution units
e.g. one for float, int, boolean
- Multiple instructions fetched and decoded at a time
- Instructions held in buffer to be executed
- Issue: no specific order to execute buffered instructions

Pipeline:

- Execute instruction n , decode $n + 1$ and fetch $n + 2$
- Any fetched instruction must be executed
- Issue: conditional statements





CPUs and cores:

- A CPU core can hold the state of more than one thread
- A core can switch between threads in a nanosecond time scale
- The OS sees several CPUs instead of one core
- Issue: what happens if there are more than two such cores?

Modern terminology¹:

- CPU: computing component (the physical entity)
- Number of cores: number of independent CPUs in a computing component
- Number of threads: maximum number of instructions that can be passed through or processed by a single core
- Number of logical cores: number of cores times number of threads

¹source: [ARK](#)

Access time

1 ns

2 ns

10 ns

10 ms

10 s

Registers

Cache

Main memory

Magnetic disk

Magnetic tape

Capacity

< 1 KB

4 MB

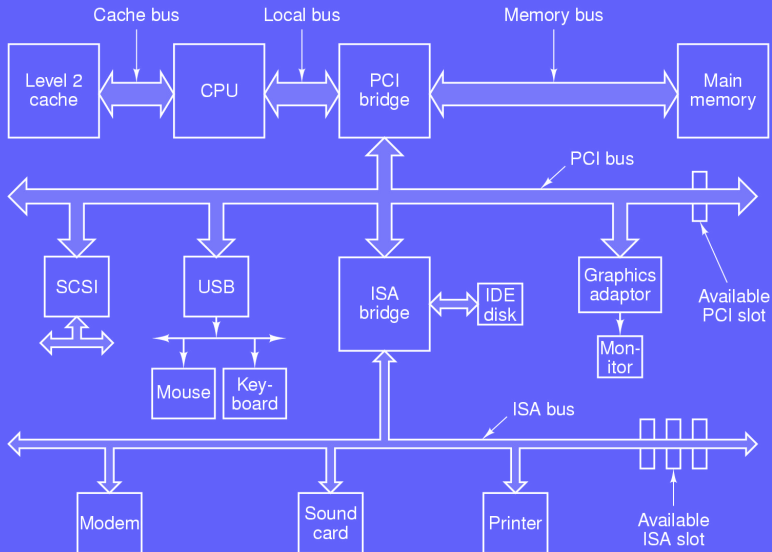
1–8 GB


200–1000 GB

400–800 GB

Memory types:

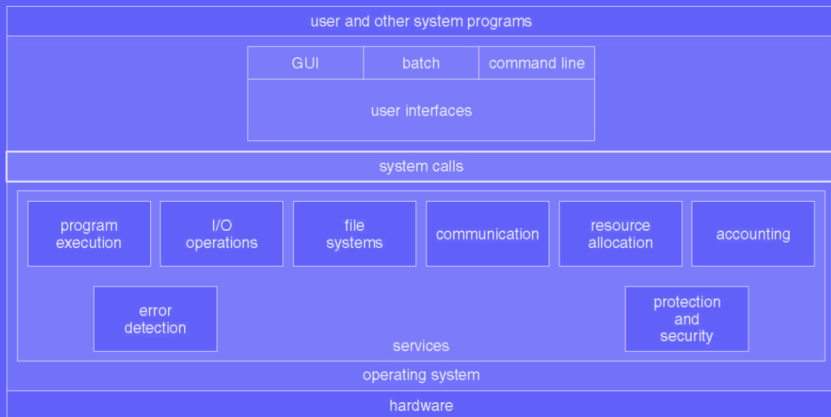
- Random Access Memory (RAM): volatile
- Read Only Memory (ROM)
- Electrically Erasable PROM (EEPROM) and flash memory: slower than RAM, non volatile.
- CMOS: save time and date , BIOS parameters
- HDD: divided into cylinder, track and sector





Five major components of an OS:

- System calls: allows to interface with user-space
- Processes: defines everything needed to run programs
- File system: store persistent data
- Input-Output (IO): allows to interface with hardware
- Protection and Security: keep the system safe



Partial list of common Unix system calls:


- Processes: `pid=fork(); pid=waitpid(pid, &statloc, options);`
`s=execve(name, argv, environp); exit(status);`
- Files: `fd=open(file,how,...); s=close(fd); s=stat(name,*buf);`
`n=read(fd,buffer,nbytes); n=write(fd,buffer,nbytes);`
`position=lseek(fd,offset,whence);`
- Directory and file system: `s=mkdir(name,mode); s=rmdir(name);`
`mount(special,name,flags,types,args); umount(name);`
`s=unlink(name); s=link(name1,name2);`
- Misc: `s=chdir(dirname); s=chmod(name,mode); sec=time(*t);`
`s=kill(pid,signal);`

A process holds all the necessary information to run a program:

- Address space belonging to the process and containing:
 - Executable program
 - Program's data
 - Program's stack
- Set of resources:
 - Registers
 - List of open files
 - Alarms
 - List of related processes
 - Any other information required by the program

The OS hides peculiarities of the disk and other IO devices

- Data stored in files grouped into directories
- The top directory is called *root* directory
- Any file can be specified using its path name
- Each process has a working directory
- Removable devices can be mounted onto the main tree
- Block files: for storage devices such as disks
- Character files: for devices accepting or outputting character streams
- Pipe: pseudo file used to connect two processes




Hardware interrupt:

- ① Send instructions to the controller
- ② The controller signals the end
- ③ Assert a pin to interrupt the CPU
- ④ Send extra information

Software interrupt:

- A call coming from userspace
- A software interrupt handler is invoked
- System call: switch to kernel mode to run privileged instruction



Hardware interrupt:

- 1 Send instructions to the controller
- 2 The controller signals the end
- 3 Assert a pin to interrupt the CPU
- 4 Send extra information

Software interrupt:

- A call coming from userspace
- A software interrupt handler is invoked
- System call: switch to kernel mode to run privileged instruction

Operating systems are almost always interrupt driven

Simplest method:

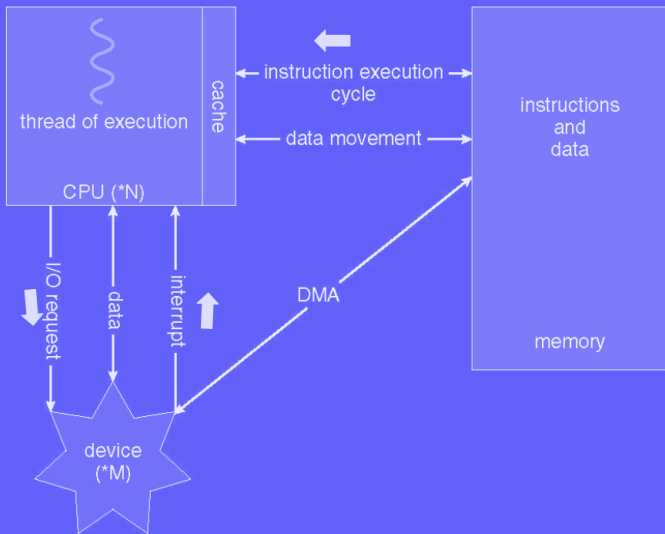
- 1 Call the driver
- 2 Start the input-output
- 3 Wait in a tight loop
- 4 Continuously poll the device to know its state

What is the drawback of this strategy?

Direct Memory Access (DMA):

- Can transmit information close to memory speeds
- Directly transfer blocks of data from the controller to the RAM
- Only little needed from the CPU
- Issue a single interrupt per block, instead of one per byte

The three communication strategies



CPU:

- Kernel Mode:
 - Set using a bit in the PSW
 - Any CPU instruction and hardware feature are available
- User mode:
 - Only a subset of instructions/features is available
 - Setting PSW kernel mode bit forbidden

Memory:


- Base and limit registers: holds the smallest legal physical memory address and the size of the range, respectively
- Memory outside the address space is protected

Input and Output:


- They are all privilege instructions
- The OS processes them to ensure their correctness and legality

Common operating system structures

Monolithic kernel



Micro kernel



- What is the main job of an OS?
- Why are there so many types of OS?
- Why is hardware important when writing an OS?
- What are the main components of an OS?
- What are system calls?



Thank you!