




Introduction to Operating Systems

2. Processes and threads


Manuel – Fall 2020





A *process* is an abstraction of a running program:

- At the core of the OS
- Process is the unit for resource management
- Oldest and most important concept
- Turn a single CPU into multiple virtual CPUs
- CPU quickly switches from process to process
- Each process run for 10-100 ms
- Processes hide the effect of interrupts

One program counter



Four program counters



Multiprogramming strategies and issues:

- CPU switches rapidly back and forth among all the processes
- Rate of computation of a process is not uniform/reproducible
- Potential issue under time constraints; e.g.
 - Read from tape
 - Idle loop for tape to get up to speed
 - Switch to another process
 - Switch back... too late

Differences between programs and processes:

- Running twice a program generates two processes
- Program: sequence of operations to perform
- Process: program, input, output, state

Differences between programs and processes:

- Running twice a program generates two processes
- Program: sequence of operations to perform
- Process: program, input, output, state

Describe the process of baking a cake when the phone rings

Times at which processes are created:


- System initialization
- Upon a user launching a new program
- Initialization of a batch job

Any created processes ends at some stage:

- Voluntarily:
 - The work is completed, issue a system call to inform the OS
 - An error is noticed, the process exits nicely
- Involuntary:
 - Fatal error, program crashes
 - Another process kills it


Two main approaches:

- UNIX-like systems:
 - A parent creates a child
 - A child can create its own child
 - The hierarchy is called *process group*
 - It is impossible to disinherit a child
- Windows system:
 - All processes are equal
 - A parent has a token to control its child
 - A token can be given to another process



Possible states:

- ① Waiting for some input
- ② Scheduler picks another process
- ③ Scheduler picks this process
- ④ Input becomes available



Possible states:

- ① Waiting for some input
- ② Scheduler picks another process
- ③ Scheduler picks this process
- ④ Input becomes available

Change of perspective on the inside of the OS:


- Do not think in terms of interrupt but of process
- Lowest level of the OS is the scheduler
- Interrupt handling, starting/stopping processes are hidden in the scheduler

A simple model for processes:

- A process is a data structure called *process control block*
- The structure contains important information such as:
 - State
 - Program counter
 - Stack pointer
 - Memory allocation
 - Open files
 - Scheduling information
- All the processes are stored in an array called *process table*

Upon an interrupt the running process must be paused:

- 1 Push on the stack the user program counter, PSW, etc.
- 2 Load information from interrupt vector
- 3 Save registers (assembly)
- 4 Setup new stack (assembly)
- 5 Finish up the work for the interrupt
- 6 Decides which process to run next
- 7 Load the new process, i.e. memory map, registers, etc. (assembly)



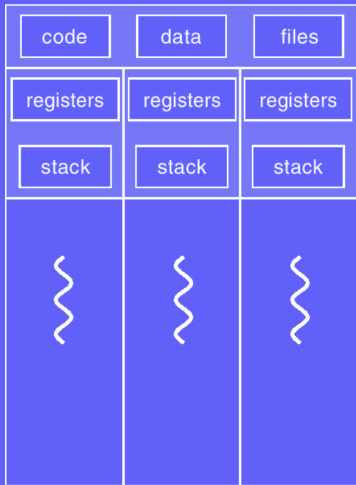

A thread is the basic unit of CPU utilisation consisting of:

- A thread ID
- The program counter
- A register set
- A stack space

All the threads within a process share the same:

- Code section
- Data section
- Operating system resources


Single vs. multi-threaded




Processes and threads:

- A thread has the same possible states as a process
- Transitions are similar to the case of a process
- Threads are sometimes called lightweight processes
- No protection is required for threads, compared to processes
- A process starts with one threads and can create more
- Processes want as much CPU as they can
- Threads can give up the CPU to let others using it

Example of a word processor





The pthread library has over 60 function calls:


- Create a thread: `int pthread_create(pthread_t *thread, const pthread_attr_t *attr, void *(*start_routine) (void *), void *arg);`
- Terminate a thread: `void pthread_exit(void *retval);`
- Wait for a specific thread to end:
`int pthread_join(pthread_t thread, void **retval);`
- Release CPU to let another thread run: `int pthread_yield(void);`
- Create and initialise a thread attribute structure:
`int pthread_attr_init(pthread_attr_t *attr);`
- Delete a thread attribute object:
`int pthread_attr_destroy(pthread_attr_t *attr);`

Creating ten threads and printing their ID

threads.c

```
1  #include <stdio.h>
2  #include <stdlib.h>
3  #include <pthread.h>
4  #define THREADS 10
5  void *gm(void *tid) {
6      printf("Good morning from thread %lu\n",*(unsigned long int*)tid);
7      pthread_exit(NULL);
8  }
9  int main () {
10     int status, i; pthread_t threads[THREADS];
11     for(i=0;i< THREADS;i++) {
12         printf("thread %d\n",i);
13         status=pthread_create(&threads[i],NULL,gm,(void*)&(threads[i]));
14         if(status!=0) {
15             fprintf(stderr,"thread %d failed with error %d\n",i,status);
16             exit(-1);
17         }
18     }
19     exit(0);
20 }
```

Threads in user-space – N:1




User-space threads:

- Kernel thinks it manages single threaded processes
- Threads implemented in a library
- Thread table similar to process table, managed by run-time system
- Switching thread does not require to trap the kernel

Questions.

- What if a thread issues a blocking system call?
- Threads within a process have to voluntarily give up the CPU




Kernel space thread:

- Kernel manages the thread table
- Kernel calls are issued to request a new thread
- Calls that might block a thread are implemented as system call
- Kernel can run another thread in the meantime

Questions.

- Why does it have a much higher cost than user space threads?
- Signals are sent to processes, which thread should received it?



Hybrid threads:

- Compromise between user-level and kernel-level
- Threading library schedules user threads on available kernel threads

Questions.

- How to implement hybrid threads?
- How to handle scheduling?

Best thread approach:

- Hybrid looked attractive
- Most systems are coming back to 1:1
- Different approaches exist on how to use threads
e.g. thread blocks on “receive system call” vs. pop up threads
- Switching implementation from single thread to multiple thread is not easy task
- Requires redesigning the whole system
- Backward compatibility must be preserved
- Research still going on to find better ways to handle threads



Thank you!