




Introduction to Operating Systems

3. Interprocess communication

Manuel – Fall 2021



In single tasking all threads are independent:

- They cannot affect or be affected by anything
- Their state is not shared with other threads
- The input state determines the output
- Everything is reproducible
- Stopping or resuming does not lead to any side effect

It suffices to run a thread to completion and start the next one

Difficulties appear with multi-tasking:

- A thread runs on one core at a time
- A thread can run on different cores at different times
- Each core is shared among several threads
- Several cores run several threads in parallel
- The number of cores has no impact on the running of the threads

Changes made by one thread can affect others

Setup for threads:

- Several threads share a common global variable
- The execution sequence impacts the global variable
- By default the behavior is random and irreproducible

Setup for threads:

- Several threads share a common global variable
- The execution sequence impacts the global variable
- By default the behavior is random and irreproducible

Major problems:

- How can threads share information?
- How to prevent them from getting on each other's way?
- How to ensure an acceptable running order?


Setup for threads:

- Several threads share a common global variable
- The execution sequence impacts the global variable
- By default the behavior is random and irreproducible

Major problems:

- How can threads share information?
- How to prevent them from getting on each other's way?
- How to ensure an acceptable running order?

All those thread issues within a process can be extended to processes within the operating system



In the printing spool:

- 1 *A* wants to queue a file: reads `next_free_slot=7`
- 2 An interrupt occurs
- 3 *B* wants to queue a file: reads `next_free_slot=7`
- 4 *B* queues its file in slot 7, and updates `next_free_slot=8`
- 5 *A* queues its file in slot 7



9:00 am

9:15 am

9:30 am

9:45 am



9:00 am

9:15 am

9:30 am

9:45 am



9:00 am

9:15 am

9:30 am

9:45 am



9:00 am

9:15 am

9:30 am

9:45 am




9:00 am

9:15 am

9:30 am

9:45 am




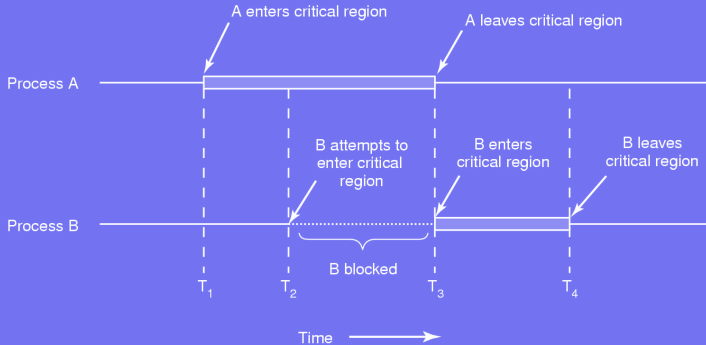
9:00 am

9:15 am

9:30 am

9:45 am





Part of the program where shared memory is accessed:

- No two processes can be in a critical region at a same time
- No assumption on the speed or number of CPUs
- No process outside a critical region can block other processes
- No process waits forever to enter a critical region

Frank

```
1  if(no milk && no note) {  
2      leave note;  
3      milk the cow;  
4      remove note;  
5  }
```

John

```
1  if(no milk && no note) {  
2      leave note;  
3      milk the cow;  
4      remove note;  
5  }
```

Frank

```
1  if(no milk && no note) {  
2      leave note;  
3      milk the cow;  
4      remove note;  
5  }
```

John

```
1  if(no milk && no note) {  
2      leave note;  
3      milk the cow;  
4      remove note;  
5  }
```

```
1  leave note Frank;  
2  if(no note John) {  
3      if(no milk) milk the cow;  
4  }  
5  remove note Frank;
```

```
1  leave note John;  
2  if(no note Frank) {  
3      if(no milk) milk the cow;  
4  }  
5  remove note John;
```

What is the issue with those two strategies?

Frank

```
1 leave note Frank;
2 while(note John) {
3     nothing;
4 }
5 if(no milk) {
6     milk the cow;
7 }
8 remove note Frank;
```

John

```
1 leave note John;
2 if(no note Frank) {
3     if(no milk) {
4         milk the cow;
5     }
6 }
7 remove note John;
```

How good is this strategy?

Symmetric strategy for two processes:

- When wanting to enter a critical region a process:
 - Shows its interest for the critical region
 - If it is accessible it exits the function and accesses it
 - If it is not accessible it waits in a tight loop
- When a process has completed its work in the critical region it signals its departure

What is the main drawback of this strategy?

Pseudo C code for two processes represented as 0 and 1

```
1  #define TRUE 1
2  #define FALSE 0
3  int turn;
4  int interested[2];
5  void enter_region(int p) {
6      int other;
7      other=1-p;
8      interested[p]=TRUE;
9      turn=p;
10     while(turn==p && interested[other]==TRUE)
11 }
12 void leave_region(int p) {
13     interested[p]=FALSE;
14 }
```

Side effects of Peterson's idea:

- Two processes: L, low priority, and H, high priority
- L enters in a critical region
- H becomes ready
- H has higher priority so the scheduler switches to H
- L has lower priority so is not rescheduled as long as H is busy
- H loops forever

Prevent the process in the critical region from being stopped:

- Disable interrupts:
 - Can be done within the kernel for a few instructions
 - Cannot be done by user processes
 - Only works when there is a single CPU
 - An interrupt on another CPU can still mess up the shared variable
- Use atomic operations:
 - Either happens in its entirety or not at all
 - Several operations can be performed at once, e.g. $A = B$
 - Requires the CPU to support the atomic update of a memory space
 - Can be used to prevent other CPUs to access a shared memory

A simple atomic operation:

- Test and Set Lock: TSL
- Copies LOCK to a register and set it to 1
- LOCK is used to coordinate the access to a shared memory
- Ensures LOCK remains unchanged while checking its value

```
1  enter_region:
2      TSL REGISTER,LOCK
3      CMP REGISTER,#0
4      JNE enter_region
5      RET
6
7  leave_region:
8      MOVE LOCK,#0
9      RET
```


A simple atomic operation:

- Test and Set Lock: TSL
- Copies LOCK to a register and set it to 1
- LOCK is used to coordinate the access to a shared memory
- Ensures LOCK remains unchanged while checking its value

```
1  enter_region:
2      TSL REGISTER,LOCK
3      CMP REGISTER,#0
4      JNE enter_region
5      RET
6
7  leave_region:
8      MOVE LOCK,#0
9      RET
```

Is this strategy better than Peterson's idea?

```
1  #define N 100
2  int count=0;
3  void producer() {
4      int item;
5      while(1) {
6          item=produce_item(); if(count==N) sleep();
7          insert_item(item); count++;
8          if(count==1) wakeup(consumer);
9      }
10 }
11 void consumer() {
12     int item;
13     while(1) {
14         if(count==0) sleep();
15         item=remove_item(); count--;
16         if(count==N-1) wakeup(producer); consume_item(item);
17     }
18 }
```

Is this code exhibiting any problem?

Assume the buffer is empty:

- Consumer reads `count == 0`
- Scheduler stops the consumer and starts the producer
- Producer adds one item
- Producer wakes up the consumer
- Consumer not yet asleep, signal is lost
- Consumer goes asleep
- When the buffer is full the producer falls asleep
- Both consumer and producer sleep forever

Basics:

- Introduced by Dijkstra in 1965
- Simple hardware based solution
- Basis of all modern OS synchronization mechanisms

Basics:

- Introduced by Dijkstra in 1965
- Simple hardware based solution
- Basis of all modern OS synchronization mechanisms

A semaphore `sem` is:

- A positive integer variable
- Only changed or tested through two actions

```
1 down(sem) {  
2     while(sem==0) sleep();  
3     sem--;  
4 }
```

```
1 up(sem) {  
2     sem++;  
3 }
```

The down operation

- If $sem > 0$, decrease it and continue
- If $sem = 0$, sleep and do not complete the down

The up operation

- Increment the value of the semaphore
- An awoken sleeping process can complete its down

The down operation

- If $sem > 0$, decrease it and continue
- If $sem = 0$, sleep and do not complete the down

The up operation

- Increment the value of the semaphore
- An awoken sleeping process can complete its down

Checking or changing the value and sleeping are done atomically:

- Single CPU: disable interrupts
- Multiple CPUs: use TSL to ensure only one CPU accesses the semaphore

Is disabling the interrupts to process the semaphore an issue?

Using semaphores to hide interrupts:

- Each IO device gets a semaphore initialised to 0
- A process accessing the device applies a `down`
- The process becomes blocked
- An interrupt is issued when the device has completed the work
- The interrupt handler processes the interrupt and applies an `up`
- The process becomes ready

A mutex is a semaphore taking values 0 (unlocked) or 1 (locked)

On a mutex-lock request:

- If the mutex is unlocked:
 - Lock the mutex
 - Enter the critical region
- If mutex is locked: put the calling thread asleep
- When the thread in the critical region exits:
 - Unlock the mutex
 - Allow a thread to acquire the lock and enter the critical region

Mutexes can be implemented in user-space using TSL

```
1 mutex-lock:
2     TSL REGISTER,MUTEX
3     CMP REGISTER,#0
4     JZ ok
5     CALL thread_yield
6     JMP mutex-lock
7 ok: RET
8
9 mutex-unlock:
10    MOVE MUTEX,#0
11    RET
```

Questions:


- What differences were introduced compared to `enter_region` (3.16)?
- In user-space what happens if a thread tries to acquire lock through busy-waiting?
- Why is `thread_yield` used?

consumer_producer.c

```
1  #include <stdio.h>
2  #include <pthread.h>
3  #define MAX 1000
4  pthread_mutex_t m; pthread_cond_t cc, cp; int buf=0;
5  void *prod() {
6      for(int i=1;i<MAX;i++) {
7          pthread_mutex_lock(&m); while(buf!=0) pthread_cond_wait(&cp,&m);
8          buf=1; pthread_cond_signal(&cc); pthread_mutex_unlock(&m);
9      }
10     pthread_exit(0);
11 }
12 void *cons() {
13     for(int i=1;i<MAX;i++) {
14         pthread_mutex_lock(&m); while(buf==0) pthread_cond_wait(&cc,&m);
15         buf=0; pthread_cond_signal(&cp); pthread_mutex_unlock(&m);
16     }
17     pthread_exit(0);
18 }
19 int main() {
20     pthread_t p, c;
21     pthread_mutex_init(&m,0); pthread_cond_init(&cc,0); pthread_cond_init(&cp,0);
22     pthread_create(&c,0,cons,0); pthread_create(&p,0,prod,0);
23     pthread_join(p,0); pthread_join(c,0);
24     pthread_cond_destroy(&cc); pthread_cond_destroy(&cp); pthread_mutex_destroy(&m);
25 }
```

Alter the previous program such as:

- To display information on the consumer and producer
- To increase the buffers to 100
- To have two consumers and one producer. In this case also print which consumer is active.



```
1  mutex mut = 0; semaphore empty = 100; semaphore full = 0;
2  void producer() {
3      while(TRUE) {
4          item = produce_item();
5          mutex-lock(&mut);
6          down(&empty); insert_item(item);
7          mutex-unlock(&mut);
8          up(&full);
9      }
10 }
11 void consumer() {
12     while(TRUE) {
13         down(&full);
14         mutex-lock(&mut); item = remove_item(); mutex-unlock(&mut);
15         up(&empty); consume_item(item);
16     }
17 }
```

Is this code working as expected?

In the previous code:

- What is the behavior of the producer when the buffer is full?
- What about the consumer?
- What is the final result for this program?
- How to fix it?


Monitors are an attempt to merge synchronization with OOP

Basic idea behind monitors:

- Programming concept that must be known by the compiler
- The mutual exclusion is not handled by the programmer
- Locking occurs automatically
- Only one process can be active within a monitor at a time
- A monitor can be seen as a “special type of class”
- Processes can be blocked and awaken based on condition variables and `wait` and `signal` functions


```
1  monitor ProducerConsumer {
2      condition full, empty;
3      int count;
4      void insert(item) {
5          if (count == N) wait(full);
6          insert_item(item);
7          count++;
8          if (count==1) signal(empty);
9      }
10     void remove() {
11         if (count==0) wait(empty);
12         removed = remove_item;
13         count--;
14         if (count==N-1) signal(full);
15     }
16     count:= 0;
17 }
```

```
1  void ProducerConsumer::producer() {
2      while (TRUE) {
3          item = produce_item();
4          ProducerConsumer.insert(item);
5      }
6  }
7  void ProducerConsumer::consumer() {
8      while (TRUE) {
9          item=ProducerConsumer.remove();
10         consume_item(item)
11     }
12 }
```



Useful when several processes must complete before the next phase



Thank you!