



Bash Pipeline

P1-pre2

Bash pipeline syntax

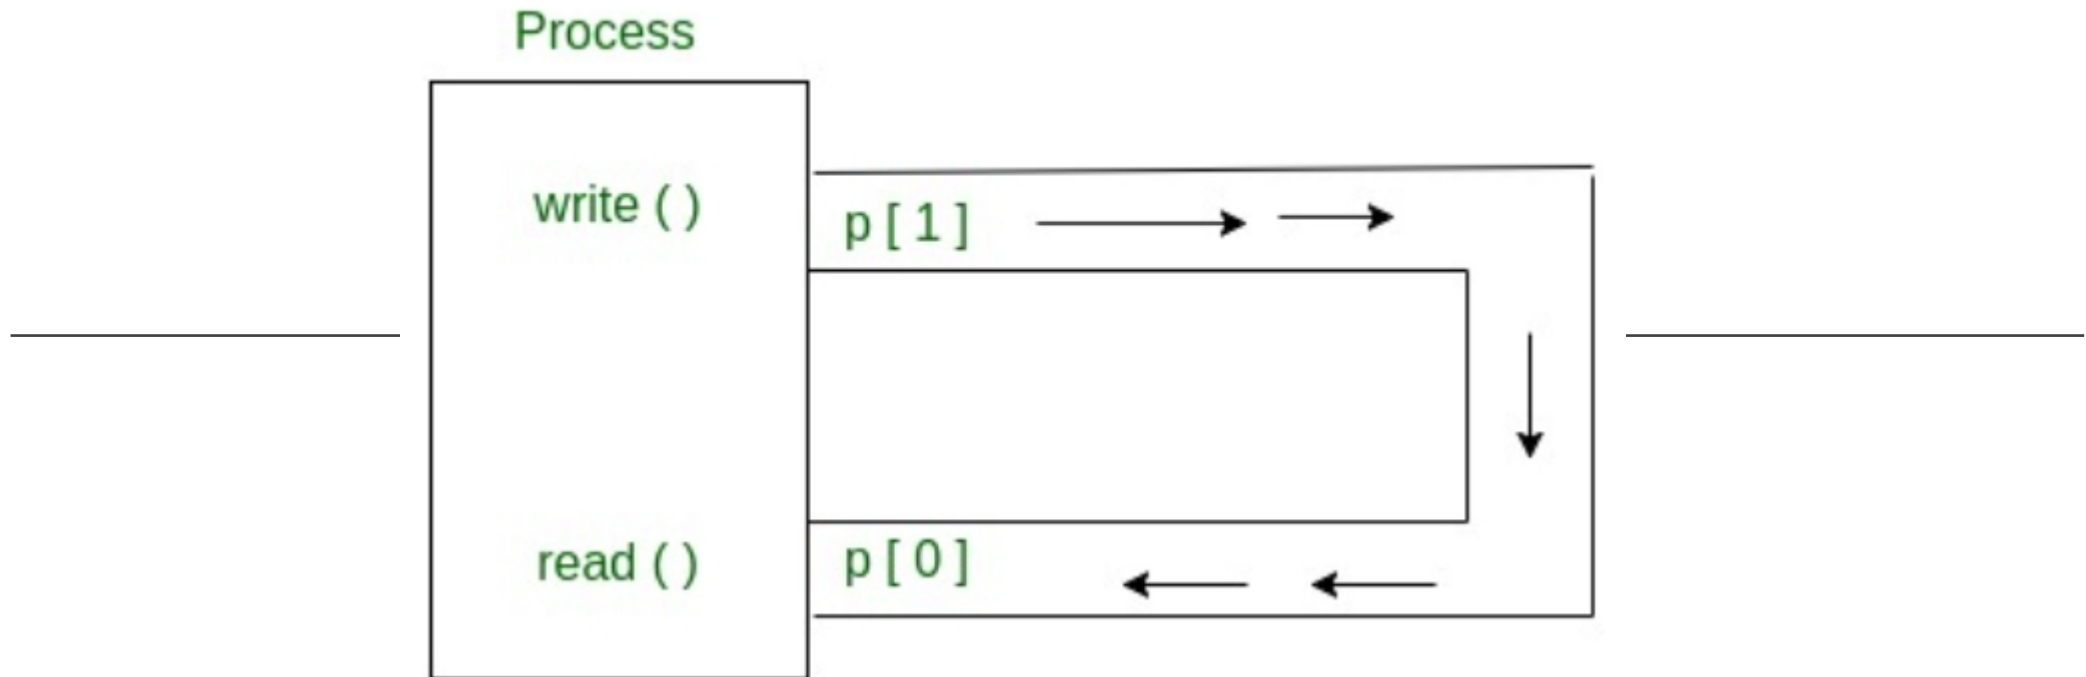
Syntax: `[command1] | [command2] | ... | [command_n]`

The pipeline “|” is used for transferring the output of the previous command to the next command’s input.

For example, “`ls | cat`”, which will connect the output of “`ls`”, which is all the files in the current directory to the “`cat`” command’s input and output to the stdout by the “`cat`” command.

Pipe() system call

- Pipe is one-way communication only, we can use a pipe such that one process write to the pipe, and the other process reads from the pipe.
- If a process tries to read before something is written to the pipe, the process is suspended until something is written.



Pipe() continued

Function prototype:

```
int pipe(int fd[2]);
```

Where `fd[0]` is the file descriptor for the read end of the pipe, `fd[1]` is the file descriptor for the write end of the pipe, like the previous picture has shown.

If a pipe is created successfully, it will return 0, if failed, it will return -1.

Pipe behave like a queue structure, which serves as FIFO (first in first out).

Pipe() continued

When we use `fork()` in any process, file descriptors remain open across child process and also parent process. If we call `fork` after creating a pipe, then the parent and child can communicate via the pipe, in other words, share the pipe.

Pipe() continued

```
#include <stdio.h>
#include <unistd.h>
#define MSGSIZE 16
char* msg1 = "hello, world #1";
char* msg2 = "hello, world #2";
char* msg3 = "hello, world #3";

int main()
{
    char inbuf[MSGSIZE];
    int p[2], i;

    if (pipe(p) < 0)
        exit(1);

    /* continued */
    /* write pipe */

    write(p[1], msg1, MSGSIZE);
    write(p[1], msg2, MSGSIZE);
    write(p[1], msg3, MSGSIZE);

    for (i = 0; i < 3; i++) {
        /* read pipe */
        read(p[0], inbuf, MSGSIZE);
        printf("%s\n", inbuf);
    }
    return 0;
}
```

The code on the left creates a pipeline with the code “pipe(p)”, then it writes the three messages defined above into the write end of the pipeline.

In the end, it takes a loop to output the buffer stored in the pipeline, which will output the three messages one by one. This is because the reading process of the pipeline can only read a little number of characters at a same time.

Thank you